

50.007 Machine Learning

Ensemble Models

Roy Ka-Wei Lee
Assistant Professor, ISTD, SUTD

Learning Objectives


You should be able to know:

- What is Generalization? Underfitting and overfitting
- Types of ensemble classifiers
 - Bagging (Random Forest)
 - Boosting
- How to model evaluation

Generalization


- Training data: $\{x_i, y_i\}$
 - Examples that we use to train our model
 - E.g. past records of days labelled with Roy goes/not go fishing
- Future/test data: $\{x_i, \hat{y}_i\}$
 - Examples that our model has not seen before
 - E.g. information about tomorrow to predict if Roy goes fishing
- Want to do well on future data not training data! $\longrightarrow$ for future unseen data.
 - Not very useful to do well on training data; we already know the label
 - Easy to be perfect on training data
 - Doing well in training does not mean will do well on future data!

Classification Example




- Find a way to classify the points

Classification Example




- Simple solution
- Unable to capture all salient pattern in data

Classification Example



- Complex solution
- “Try too hard” to work well with training data
- Fit noise into the model
- Pattern is one-off and may not appear again in future data


Classification Example




- Good solution, able to generalize well for future data

but do not overfit !


Underfitting and Overfitting



Underfitting
(Too simple to
explain the
variance)



Overfitting
(Too complex
- try too hard)



Appropriate
Fitting ✨

Underfitting and Overfitting

• Overfitting

- Model is too complex (too flexible)
- Fit noise in training data
- Pattern is one-off and might not appear in future data
- Model F overfits the data if:

- We can find another model F'
- Which makes more mistakes in training data: $E_{\text{Train}}(F') > E_{\text{Train}}(F)$
- But less mistake in future data: $E_{\text{Test}}(F) > E_{\text{Test}}(F')$


• Underfitting

- Model is too simple
- Not powerful enough to capture the salient pattern
- Can find another model F' with smaller $E_{\text{Test}}(F')$ and $E_{\text{Train}}(F')$

intuition


Error in model F' on training
✓
Error in model F on training

Underfitting and Overfitting





As you ↑ the # of nodes,
the % error ↑ as well.

Underfitting and Overfitting



What Causes Overfitting?

- Noise in training data
- * • Insufficient data points in training data



Further Notes on Overfitting

- Overfitting results in decision trees that are more complex than necessary
- Training error no longer provides a good estimate of how well the tree will perform on previously unseen records
- Need new ways for estimating errors

Estimating Generalization Errors ★ IMPORTANT MCQ / T&F

- Training errors: error on training ($\sum e(t)$)
- Generalization errors: error on testing data ($\sum e'(t)$)
- Methods for **estimating** generalization errors:
 - ① ○ **Optimistic approach**: simply use training error
 - $e'(t) = e(t)$ $\Rightarrow$ Error on training = Error on testing
 - ② ○ **Pessimistic approach**:
 - For each leaf node: $e'(t) = (e(t) + 0.5)$ $\rightarrow$ arbitrary / estimated.
 - Total errors: $e'(T) = e(T) + N \times 0.5$ (N: number of leaf nodes)
 - Example: For a tree with 30 leaf nodes and 10 errors on training data (out of 1000 training data instances):
 - ① • Training error = $10/1000 = 1\%$
 - ② • Generalization error = $(10 + 30 \times 0.5)/1000 = 2.5\%$ $\rightarrow$ Estimation
 - **Reduced error pruning (REP)**:
 - uses validation data set to estimate generalization error



Estimating Generalization Errors

Training

S/N	A	B	C	Label
1	0	0	0	+
2	0	0	1	+
3	0	1	0	+
4	0	1	1	-
5	1	0	0	+


✓
✓
✗
✓
✗

Testing

S/N	A	B	C	Label
6	0	0	0	+
7	0	1	1	+
8	1	1	0	+
9	1	0	1	-
10	1	0	0	+

✓
✗
✗
✗
✗

Model



Training Error = $\frac{2}{5} = 0.4$

Pessimistic Error:
 $= (2 + 4 \times 0.5) / 5$
 $= 0.8$

Actual = $\sum e'(t)$
 $= 4/5 = 0.8$

Estimating Generalization Error

Optimistic: ? $e(t) : e'(t) = \frac{2}{5} = 0.4$

Pessimistic: ? $e'(t) = e(t) + N \times 0.5$
 $= (2 + 4 \times 0.5) / 5 = 0.8$

Actual Generalization Error : ?

$\sum e'(t) = \frac{4}{5} = 0.8$


Estimating Generalization Errors

Training

S/N	A	B	C	Label	
1	0	0	0	+	✓
2	0	0	1	+	✓
3	0	1	0	+	✗
4	0	1	1	-	✓
5	1	0	0	+	✗

Testing

S/N	A	B	C	Label	
6	0	0	0	+	✓
7	0	1	1	+	✗
8	1	1	0	+	✗
9	1	0	1	-	✗
10	1	0	0	+	✗



Estimating Generalization Error

Optimistic: ?

Pessimistic: ?

Actual Generalization Error : ?


Estimating Generalization Errors

Training

S/N	A	B	C	Label	
1	0	0	0	+	✓
2	0	0	1	+	✓
3	0	1	0	+	✗
4	0	1	1	-	✓
5	1	0	0	+	✗

Testing

S/N	A	B	C	Label	
6	0	0	0	+	✓
7	0	1	1	+	✗
8	1	1	0	+	✗
9	1	0	1	-	✗
10	1	0	0	+	✗



Estimating Generalization Error *normalising*
 Optimistic: $e'(t) = e(t) = 2 / 5 = 0.4$

Pessimistic: $e'(t) = (e(t) + 0.5) = (2 + 4 * 0.5) / 5 = 0.8$

Actual Generalization Error : $e'(t) = 4 / 5 = 0.8$

Occam's Razor

- Given two models of similar generalization errors, one should prefer the simpler model over the more complex model
- For complex models, there is a greater chance that it was fitted accidentally by errors in data
- Therefore, one should include model complexity when evaluating a model

↓
Big-O complexity

How to Address Overfitting

- **Pre-Pruning (Early Stopping Rule)**


- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
 - Stop if all instances belong to the same class
 - Stop if all the attribute values are the same
 - More restrictive conditions: E.g., Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

- **Post-Pruning**


- Grow decision tree to its entirety
- Trim the nodes of the decision tree in a bottom-up fashion
 - If generalization improves after trimming, replace sub-tree by a leaf node
- Class label of leaf node is determined from majority class of instances in the sub-tree

performance
of test
improve

How will lazy Roy do it?



How will lazy Roy do it?




First split the
dataset into 3
parts...



↓
defined by user

How will lazy Roy do it?




We train the model
using the training
data and test it on
validation set...

```
DT = DecisionTreeClassifier(depth=d)  
DT.fit(training data)  
Pred = DT.predict(validation data)  
Acc = ComputeAcc(Pred)
```



How will lazy Roy do it?




We loop and try
different depth and
get the best
accuracy

```
For d in range(101,1,-1)
    DT = DecisionTreeClassifier(depth=d)
    DT.fit(training data)
    Pred = DT.predict(validation data)
    Acc = ComputeAcc(Pred)
    if BestAcc < Acc:
        BestAcc = Acc
        BestDepth = d
```



How will lazy Roy do it?



Training

Validation

Testing

After we get the best depth, we re-train the model with it and test it on testing data


```
For d in range(101,1,-1)
    DT = DecisionTreeClassifier(depth=d)
    DT.fit(training data)
    Pred = DT.predict(validation data)
    Acc = ComputeAcc(Pred)
    if BestAcc < Acc:
        BestAcc = Acc
        BestDepth = d
DT = DecisionTreeClassifier(depth=BestDepth)
DT.fit(training data)
Pred = DT.predict(testing data)
Acc = ComputeAcc(Pred)
```

→ training + validation



How will lazy Roy do it?

* For project, might need to cut from validation to test :: testing data have no labels.



Caveat: workable solution but still have some drawbacks...



It will take too slow!!!


```
For d in range(101,1,-1)
    DT = DecisionTreeClassifier(depth=d)
    DT.fit(training data)
    Pred = DT.predict(validation data)
    Acc = ComputeAcc(Pred)
    if BestAcc < Acc:
        BestAcc = Acc
        BestDepth = d
DT = DecisionTreeClassifier(depth=BestDepth)
DT.fit(training data)
Pred = DT.predict(testing data)
Acc = ComputeAcc(Pred)
```

Ensemble Methods *★ IMPORTANT*

- Construct a set of classifiers from the training data
- Predict class label of previously unseen records by aggregating predictions made by multiple classifiers
- Assumption:
 - Individual classifiers (voters) could be lousy (stupid), but the aggregate (electorate) can usually classify (decide) correctly.



General Idea



$t = \text{usually } 10 \#$
to ensure somebody wins.

Why does it work?

- Suppose there are 25 base classifiers
 - Each classifier has error rate, $\varepsilon = 0.35$
 - Assume classifiers are independent
 - Probability that the ensemble classifier makes a wrong prediction (i.e., 13 out of the 25 classifiers misclassified):

$$\sum_{i=13}^{25} \binom{25}{i} \varepsilon^i (1 - \varepsilon)^{25-i} = 0.06$$

Ensemble Methods

- How to generate an ensemble of classifiers?
(list below is non exhaustive!)
 - Bagging
 - Boosting

Bagging

- Bootstrap **A**ggregating (**B**agging)
- Multiple **models** of same learning algorithm trained with subset of dataset **randomly sampled (with replacement)** from the training dataset
- Each sample has probability $1 - (1 - 1/n)^n$ of being selected
 - This value tends to be **0.63** for large n

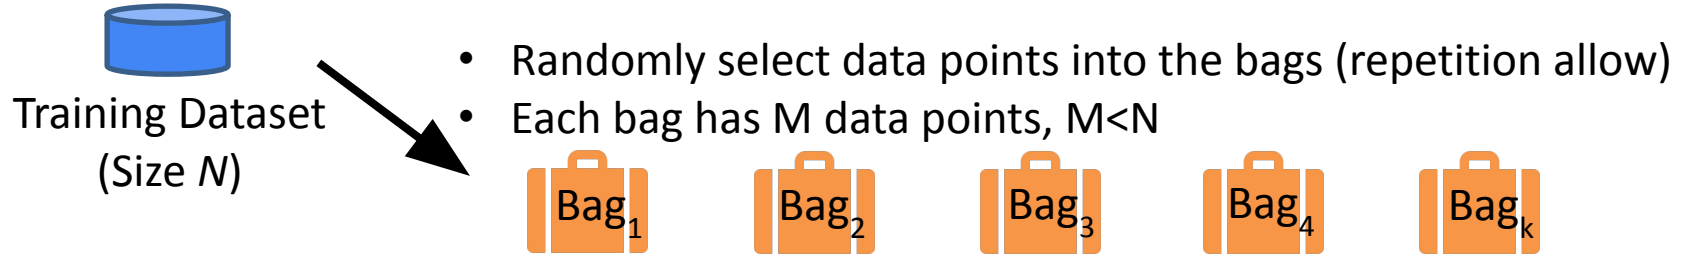
↓
due to the central limit theorem (CLT)
∴ if n is large, the P falls to a bell curve,
there is a $\uparrow P \approx 0.63$

Bagging




Training Dataset
(Size N)


Bagging



Bagging



Bagging ~~IMPORTANT~~



Random Forest

- Train a lot of decision trees and use bagging
- Wisdom of the crowd!
 - Uncorrelated trees are preferred

Random Forest

- Train a lot of decision trees and use bagging
- Wisdom of the crowd!
 - Uncorrelated trees are preferred



Random Forest

- Bootstrapping + Feature Randomness



Boosting

⇒ Analogy: For labels they guessed wrong, the model will go back & keep guessing till it is correct.

- An **iterative procedure** to adaptively change distribution of training data by focusing more on previously **misclassified records**
 - Initially, all N records are assigned **equal weights**
 - Unlike bagging, **sampling weights may change** at the end of boosting round
 - Records that are **wrongly** classified will have their **weights increased**
 - Records that are **correctly** classified will have their **weights decreased**
 - **Caveat:** boosting show **better predictive accuracy** than bagging but also **tends to over-fits the training data**

Boosting




Training Dataset
(Size N)

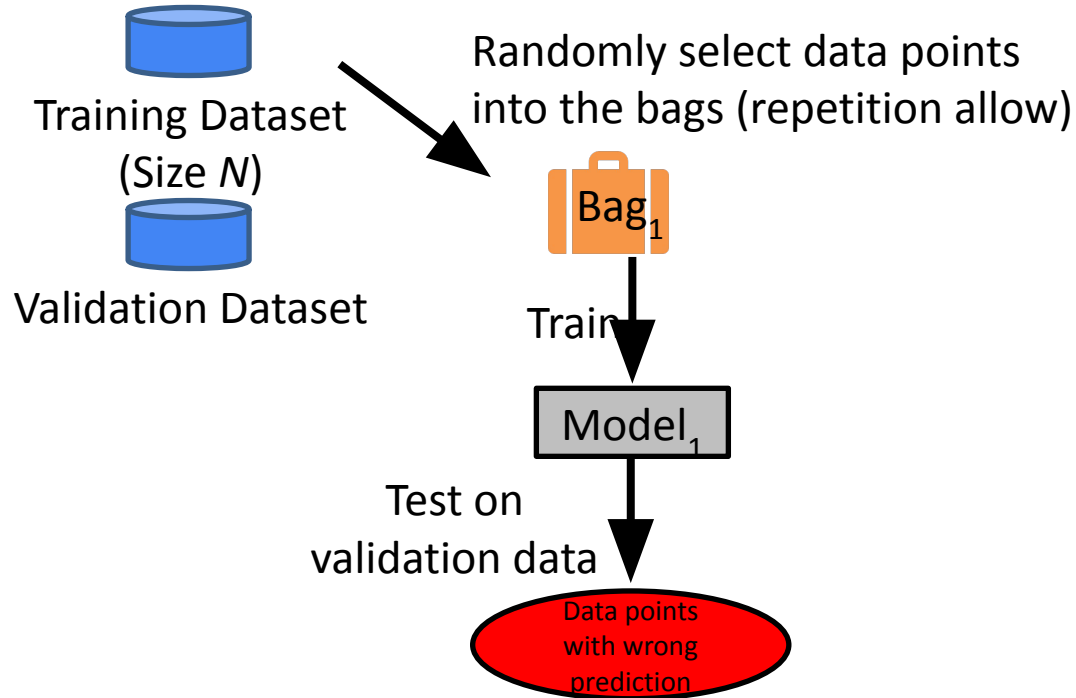


Validation Dataset

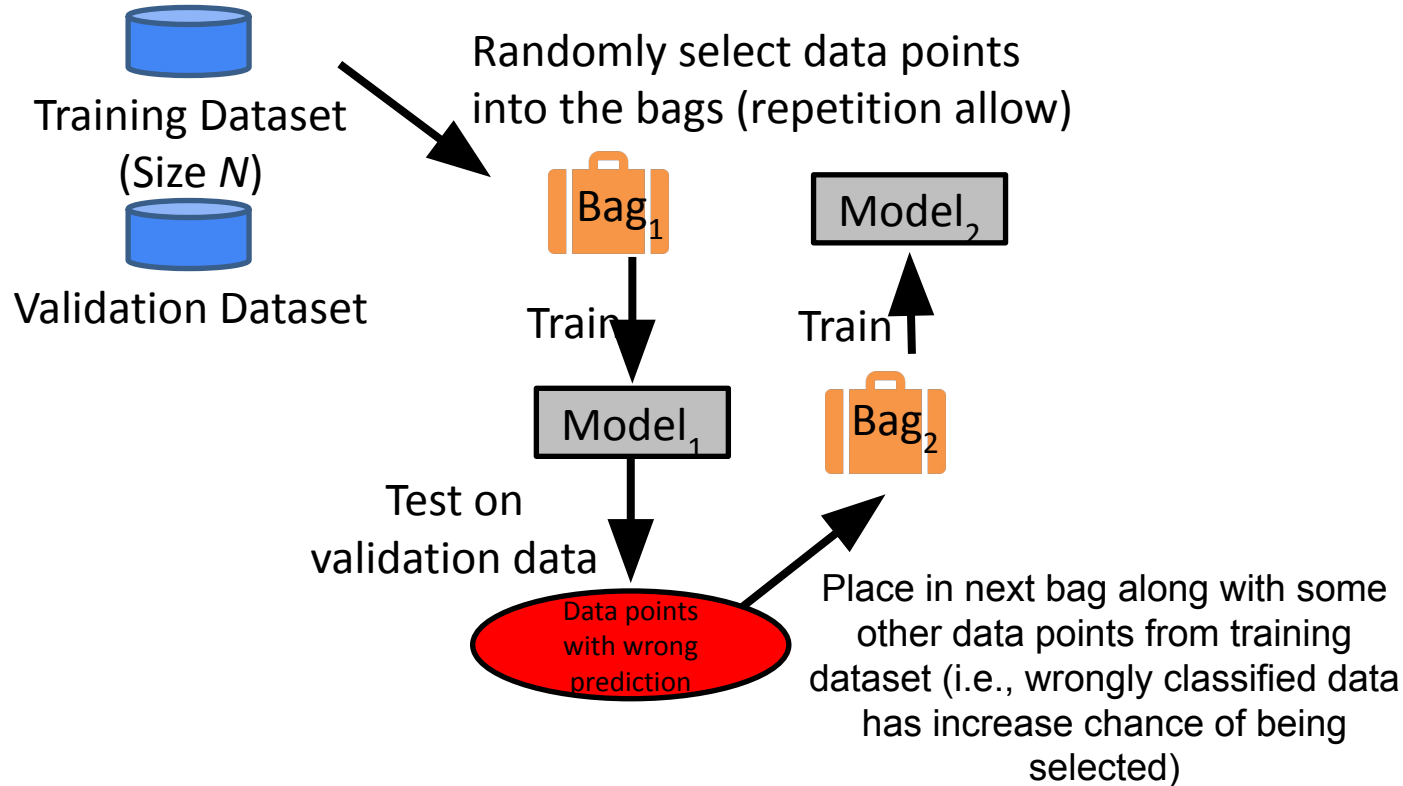
Boosting




Boosting




Boosting



Boosting




Boosting *IMPORTANT*




AdaBoost

Iteration 1




AdaBoost

Iteration 1




Iteration 2




AdaBoost


Iteration 1



Iteration 2




Iteration 3




AdaBoost


Iteration 1




Iteration 2



Iteration 3



Combine the
3 model



Final Classifier/
Strong Classifier

AdaBoost Algorithm

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: $C_1, C_2, \dots, C_T$

↓
equal chance

AdaBoost Algorithm

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: $C_1, C_2, \dots, C_T$
- For $i=1$ to T :
 - Fit a classifier $C_i(x)$ to training data
 - Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

AdaBoost Algorithm


- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: $C_1, C_2, \dots, C_T$
- For $i=1$ to T :

- Fit a classifier $C_i(x)$ to training data
- Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

- Compute importance of classifier C_i :

$$\alpha_i = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_i}{\varepsilon_i} \right)$$



When the error rate goes beyond 0.5, we assign the classifier a negative weight (i.e., we don't want them!)

$\hookrightarrow > 0.5$, worse than random, $\therefore$ importance < 0

AdaBoost Algorithm

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: $C_1, C_2, \dots, C_T$
- For $i=1$ to T :

- Fit a classifier $C_i(x)$ to training data
- Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

- Compute importance of classifier C_i :

$$\alpha_i = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_i}{\varepsilon_i} \right)$$

- Update the weights

$$w_i^{(j+1)} = \frac{w_i^{(j)}}{Z_j} \begin{cases} \exp^{-\alpha_j} & \text{if } C_j(x_i) = y_i \\ \exp^{\alpha_j} & \text{if } C_j(x_i) \neq y_i \end{cases}$$

where Z_j is the normalization factor



- $\exp^{-\alpha}$ is always lesser than 1 when the data point is correctly classified, thus giving it a smaller weight
- $\exp^{\alpha}$ is bigger than 1 when it is misclassified, assigning this data point a greater weight

AdaBoost Algorithm *★IMPORTANT*

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: $C_1, C_2, \dots, C_T$
- For $i=1$ to T :
 - Fit a classifier $C_i(x)$ to training data

- Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

- Compute importance of classifier C_i :

$$\alpha_i = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_i}{\varepsilon_i} \right)$$

- Update the weights

$$w_i^{(j+1)} = \frac{w_i^{(j)}}{Z_j} \begin{cases} \exp^{-\alpha_j} & \text{if } C_j(x_i) = y_i \\ \exp^{\alpha_j} & \text{if } C_j(x_i) \neq y_i \end{cases}$$

where Z_j is the normalization factor

- Final Classifier:

$$C^*(x) = \arg \max_y \sum_{j=1}^T \alpha_j \delta(C_j(x) = y)$$



Model Evaluation

- Metrics for Performance Evaluation
 - How to evaluate the performance of a model?
- Methods for Performance Evaluation
 - How to obtain reliable estimates?
- Methods for Model Comparison
 - How to compare the relative performance among competing models?

Metrics for Performance Evaluation

- Focus on the **predictive capability** of a model
 - Rather than how fast it takes to classify or build models, scalability, etc.
- Confusion Matrix:

	PREDICTED CLASS $\hat{y} = y'$		
y ACTUAL CLASS		Class=Yes	Class=No
	Class=Yes	a	b
	Class=No	c	d

a: TP (true positive)
b: FN (false negative)
c: FP (false positive)
d: TN (true negative)

Metrics for Performance Evaluation

ACTUAL CLASS	PREDICTED CLASS		
		Class=Yes	Class=No
	Class=Yes	a (TP)	b (FN)
	Class=No	c (FP)	d (TN)

Most widely-used metric:

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$

Limitation of Accuracy

- Consider a 2-class problem
 - Number of Class 0 examples = 9990
 - Number of Class 1 examples = 10
- If model predicts everything to be class 0, accuracy is $9990/10000 = 99.9\%$
 - Accuracy is misleading because model does not detect any class 1 example

Cost Matrix

**IMP-RTANT*

MCQ / T&F

- $C(i|j)$: Cost of misclassifying class j example as class i

ACTUAL CLASS	PREDICTED CLASS		
	$C(i j)$	Class=Yes	Class=No
	Class=Yes	$C(\text{Yes} \text{Yes})$ TP	$C(\text{No} \text{Yes})$ FN
	Class=No	$C(\text{Yes} \text{No})$ FP	$C(\text{No} \text{No})$ TN



TP: True Positive
TN: True Negative

FN: False Negative
FP: False Positive

Computing Cost of Classification

Cost Matrix	PREDICTED CLASS		
ACTUAL CLASS	C(i j)	+	-
	+	-1	100
	-	1	0

Model M_1	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	150	40
	-	60	250

$$\text{Accuracy} = ? \quad \frac{TP + TN}{\text{Total}} = \frac{150 + 250}{150 + 60 + 40 + 250}$$

$$\text{Cost} = ? \quad = \frac{400}{500} = \frac{4}{5} = 0.8$$

Model M_2	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	250	45
	-	5	200

$$\text{Accuracy} = ? \quad \frac{250 + 200}{250 + 5 + 45 + 200} = \frac{450}{500} = \frac{45}{50}$$

$$\text{Cost} = ? \quad = \frac{9}{10} = 0.9$$

$$\begin{aligned}
 \text{Cost} &= 15(-1) + 40(100) + 6(1) + 250(0) \\
 &= -150 + 4000 + 60 + 0 \\
 &= 3850 + 60 \\
 &= 3910
 \end{aligned}$$

$$\begin{aligned}
 \text{Cost} &= 250(-1) + 45(100) + 5(1) + 20(0) \\
 &= -250 + 4500 + 5 + 0 \\
 &= 4250 + 5 \\
 &= 4255
 \end{aligned}$$

$\therefore M_2$ might be have better accuracy but it has a higher cost.

M_1 might have a lower accuracy but much significant cost, $\therefore$ business might go with M_1 $\because$ of lower cost.

Computing Cost of Classification

Cost Matrix	PREDICTED CLASS		
ACTUAL CLASS	C(i j)	+	-
	+	-1	100
	-	1	0

Model M_1	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	150	40
	-	60	250

$$\text{Accuracy} = 150 + 250 / 150 + 250 + 40 + 60$$

$$= 80\%$$

$$\text{Cost} = 150(-1) + 40(100) + 60(1) + 250(0)$$

$$= 3910$$

Model M_2	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	250	45
	-	5	200

$$\text{Accuracy} = 250 + 200 / 250 + 200 + 45 + 5$$

$$= 90\%$$

$$\text{Cost} = 250(-1) + 45(100) + 5(1) + 200(0)$$

$$= 4255$$

Precision, Recall & F1 ~~★~~ IMPRIANT

	PREDICTED CLASS		
		Class=Yes	Class=No
	ACTUAL CLASS		
	Class=Yes	a (TP)	b (FN)
	Class=No	c (FP)	d (TN)

$$\text{Precision (p)} = \frac{a}{a + c}$$

$$\frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall (r)} = \frac{a}{a + b}$$

$$\frac{\text{TP}}{\text{TP} + \text{FN}}$$


$$\text{F - measure (F)} = \frac{2rp}{r + p} = \frac{2a}{2a + b + c}$$



Methods for Performance Evaluation

- How to obtain a reliable estimate of performance?
- Performance of a model may depend on other factors besides the learning algorithm:
 - Class distribution
 - Cost of misclassification
 - Size of training and test sets

Learning Curve



- Learning curve shows how accuracy changes with varying sample size
- Requires a sampling schedule for creating learning curve:
 - Arithmetic sampling (Langley, et al.)
 - Geometric sampling (Provost et al.)
- Effect of small sample size:
 - Bias in the estimate
 - Variance of estimate

Methods of Estimation


- **Holdout** = Validation $\Rightarrow 80/20$
 - Reserve 2/3 for training and 1/3 for test
- **Cross validation**
 - Partition data into k disjoint subsets
 - k-fold: train on k-1 partitions, test on the remaining one
 - Leave-one-out: k=n
- **Random subsampling**
 - Repeated holdout

Train, Test and Validate

- **Training Set:** The actual dataset that we use to train the model (weights and biases in the case of Neural Network). The model sees and learns from this data.
- **Validation Set:** The sample of data used to provide an unbiased evaluation of a model fit on the training dataset while tuning model hyperparameters. The evaluation becomes more biased as skill on the validation dataset is incorporated into the model configuration.
- **Testing Set:** The sample of data used to provide an unbiased evaluation of a final model fit on the training dataset.


Cross Validation

- K-fold cross-validation
 - Create K-fold partition of the dataset.
 - Form K hold-out predictors, each time using one partition as validation and rest K-1 as training datasets.
 - Final predictor is average/majority vote over the K hold-out estimates.




Cross Validation

- Leave-one-out (LOO) cross-validation
 - Special case of K-fold with $K=n$ partitions
 - Equivalently, train on $n-1$ samples and validate on only one sample per run for n runs



Random Subsampling

- Randomly subsample a fixed fraction α ($0 < \alpha < 1$) of the dataset for validation.
- Train the model with remaining data as training data.
- Repeat K times
- Compute the average metrics for the K hold-out estimates.



What you should know?

- What is Generalization? Underfitting and overfitting
- Types of ensemble classifiers
 - Bagging (Random Forest)
 - Boosting
- How to model evaluation

Suggestion: If you're not very clear or want to learn more, please do read: "C. Bishop: Pattern Recognition and Machine Learning. Springer, 2006" (recommended text book). It will help you to understand the concept.

Acknowledgement

- *The following material have been referenced or partially used:*
 - *Lecture notes from Prof. Berrak Sisman and Prof. Zhao Na*
 - *National University of Singapore – Pattern Recognition Module (EE5907R)*
 - *National University of Singapore – Neural Networks Module (EE5904R)*
 - *University of Edinburgh, United Kingdom – Machine Learning And Pattern Recognition (MLPR, INFR11130)*
 - *Stanford University – Machine Learning (CS229)*